CURRICULUM VITAE

Bargav Jayaraman

PhD, Computer Science
Postdoctoral Researcher, Meta
1 Hacker Way, Menlo Park CA 94025

Email: bargavjayaraman@gmail.com
Web: http://bargavjayaraman.github.io/

(a) Research Interests

Privacy preserving machine learning, differential privacy, evaluation metrics for privacy leakage quantification, memorization in LLMs and VLMs, federated learning.

(b) Education & Training

2016 – 2022	Computer Science, PhD	University of Virginia	Virginia, USA
2012 – 2015	Computer Science, MS	IIIT Hyderabad	Telangana, India
2008 – 2012	Computer Science, B. Tech	SASTRA University	Tamil Nadu, India

(c) Research & Professional Experience

Jan 2023 – Now	Postdoctoral Researcher, Meta FAIR, Menlo Park
May 2021 – Aug 2021	Research Intern, Microsoft Research, Seattle
Aug 2016 – Dec 2022	Graduate Research Assistant, University of Virginia, Charlottesville
Jan 2015 – Jun 2016	R & D Senior Analyst, Accenture Technology Labs, Bangalore
Fall 2014	Teaching Assistant (DWDM), IIIT Hyderabad, Hyderabad
Spring 2014	Teaching Assistant (POIS), IIIT Hyderabad, Hyderabad

(d) Skill Set

Programming	Python, C, C++, Java
Frameworks	Scikit-Learn, PyTorch, Tensorflow, Obliv-C, AWS, Azure, Git
Web	HTML, CSS, Markdown

(e) Publications

Attacks On Machine Learning

1. Narine Kokhlikyan*, Bargav Jayaraman*, Florian Bordes, Chuan Guo, and Kamalika Chaudhuri. Measuring déjà vu memorization efficiently. Under review, 2024.
2. Bargav Jayaraman, Chuan Guo, and Kamalika Chaudhuri. Déjà vu memorization in vision-language models. *arXiv:2402.02103*, 2024.
3. Bargav Jayaraman, Esha Ghosh, Melissa Chase, Sambuddha Roy, Wei Dai, and David Evans. Combing for credentials: Active pattern extraction from smart reply. In *IEEE Symposium on Security and Privacy*, 2024.
4. Bargav Jayaraman and David Evans. Are attribute inference attacks just imputation? In *ACM Conference on Computer and Communications Security*, 2022.
5. Bargav Jayaraman, Lingxiao Wang, Katherine Knipmeyer, Quanquan Gu, and David Evans. Revisiting membership inference under realistic assumptions. In *Proceedings on Privacy Enhancing Technologies*, 2021.
6. Bargav Jayaraman and David Evans. Evaluating differentially private machine learning in practice. In *USENIX Security Symposium*, 2019.

Privacy-Preserving Machine Learning

7. Lingxiao Wang, Bargav Jayaraman, David Evans, and Quanquan Gu. Efficient privacy-preserving stochastic nonconvex optimization. In *Uncertainty in Artificial Intelligence*, 2023.
8. Bargav Jayaraman, Lingxiao Wang, David Evans, and Quanquan Gu. Distributed learning without distress: Privacy-preserving empirical risk minimization. In *Advances in Neural Information Processing Systems*, 2018.
9. Lu Tian*, Bargav Jayaraman*, Quanquan Gu, and David Evans. Aggregating private sparse learning models using multi-party computation. In *NeurIPS Workshop on Private Multi-Party Machine Learning*, 2016.

Other Publications

10. Bargav Jayaraman*, Hannah Li*, and David Evans. Decentralized certificate authorities. *arXiv:1706.03370*, 2017.
11. Breno D. Cruz, Bargav Jayaraman, Anurag Dwarakanath, and Collin McMillan. Detecting vague words and phrases in requirements documents in a multilingual environment. In *IEEE International Requirements Engineering Conference (RE)*, 2017.
12. Bruhadeshwar Bezawada, Alex X. Liu, Bargav Jayaraman, Ann L. Wang, and R. Li. Privacy preserving string matching for cloud computing. In *IEEE International Conference on Distributed Computing Systems*, 2015.

(f) Invited Talks and Presentations

1. Presented my published work on "Combing for Credentials: Active Pattern Extraction from Smart Reply" at *IEEE S & P*, 2024.
2. Presented my published work on "Are Attribute Inference Attacks Just Imputation?" at *ACM CCS*, 2022.
3. Presented my published work on "Revisiting Membership Inference Under Realistic Assumptions" at *PETS Symposium*, 2021.
4. Presented a poster on "Revisiting Membership Inference Under Realistic Assumptions" at *TPDP* and *PPML workshops* co-located with *CCS* 2020 and *NeurIPS* 2020 conferences resp.
5. Gave an invited talk at Microsoft Research in Summer 2020 where I presented my work on membership inference attacks on machine learning models.
6. Gave a talk on evaluating privacy-utility trade-off of privacy preserving machine learning at *AIML seminar* held at University of Virginia in Fall 2019.
7. Gave a talk on evaluating privacy preserving-machine learning at Winter 2019 *DCAPS workshop* held at University of Maryland College Park.
8. Presented my published work on "Evaluating Differentially Private Machine Learning in Practice" at *USENIX Security Symposium*, 2019.
9. Presented my published work on "Distributed Learning without Distress: Privacy-Preserving Empirical Risk Minimization" at *NeurIPS conference*, 2018.
10. Presented my published work on "Aggregating Private Sparse Learning Models using Multi-Party Computation" at *PPML workshop* co-located with *NeurIPS conference*, 2016.

(g) Awards and Achievements

1. Awarded travel grant at *USENIX Security Symposium*, 2019.
2. Awarded travel grant at *NeurIPS conference*, 2018.

3. Filed *three* patents while working at Accenture Technology Labs Bangalore.

(h) Professional Services

1. Program Committee Member for: NeurIPS 2024, ICML 2024, NeurIPS 2023, USENIX Security 2023, PPML Workshop 2021 and Privacy in ML Workshop 2021.
2. Reviewer for ACM CCS 2021 and IEEE TDSC 2021.